

```
In [58]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import warnings
warnings.filterwarnings("ignore")
import os
```

```
In [59]: data=pd.read_csv("car data.csv")
data.head()
```

```
Out[59]:
```

| | Car_Name | Year | Selling_Price | Present_Price | Driven_kms | Fuel_Type | Selling_type | Transmission | Owner |
|---|----------|------|---------------|---------------|------------|-----------|--------------|--------------|-------|
| 0 | ritz | 2014 | 3.35 | 5.59 | 27000 | Petrol | Dealer | Manual | 0 |
| 1 | sx4 | 2013 | 4.75 | 9.54 | 43000 | Diesel | Dealer | Manual | 0 |
| 2 | ciaz | 2017 | 7.25 | 9.85 | 6900 | Petrol | Dealer | Manual | 0 |
| 3 | wagon r | 2011 | 2.85 | 4.15 | 5200 | Petrol | Dealer | Manual | 0 |
| 4 | swift | 2014 | 4.60 | 6.87 | 42450 | Diesel | Dealer | Manual | 0 |

```
In [60]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 301 entries, 0 to 300
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Car_Name        301 non-null    object
1   Year            301 non-null    int64
2   Selling_Price   301 non-null    float64
3   Present_Price   301 non-null    float64
4   Driven_kms      301 non-null    int64
5   Fuel_Type       301 non-null    object
6   Selling_type    301 non-null    object
7   Transmission    301 non-null    object
8   Owner           301 non-null    int64
dtypes: float64(2), int64(3), object(4)
memory usage: 21.3+ KB
```

```
In [61]: data.isna().any()
```

```
Out[61]: Car_Name        False
Year            False
Selling_Price   False
Present_Price   False
Driven_kms      False
Fuel_Type       False
Selling_type    False
Transmission    False
Owner           False
dtype: bool
```

```
In [62]: print(data.Fuel_Type.value_counts(),"\n")
print(data.Selling_type.value_counts(),"\n")
print(data.Transmission.value_counts())
```

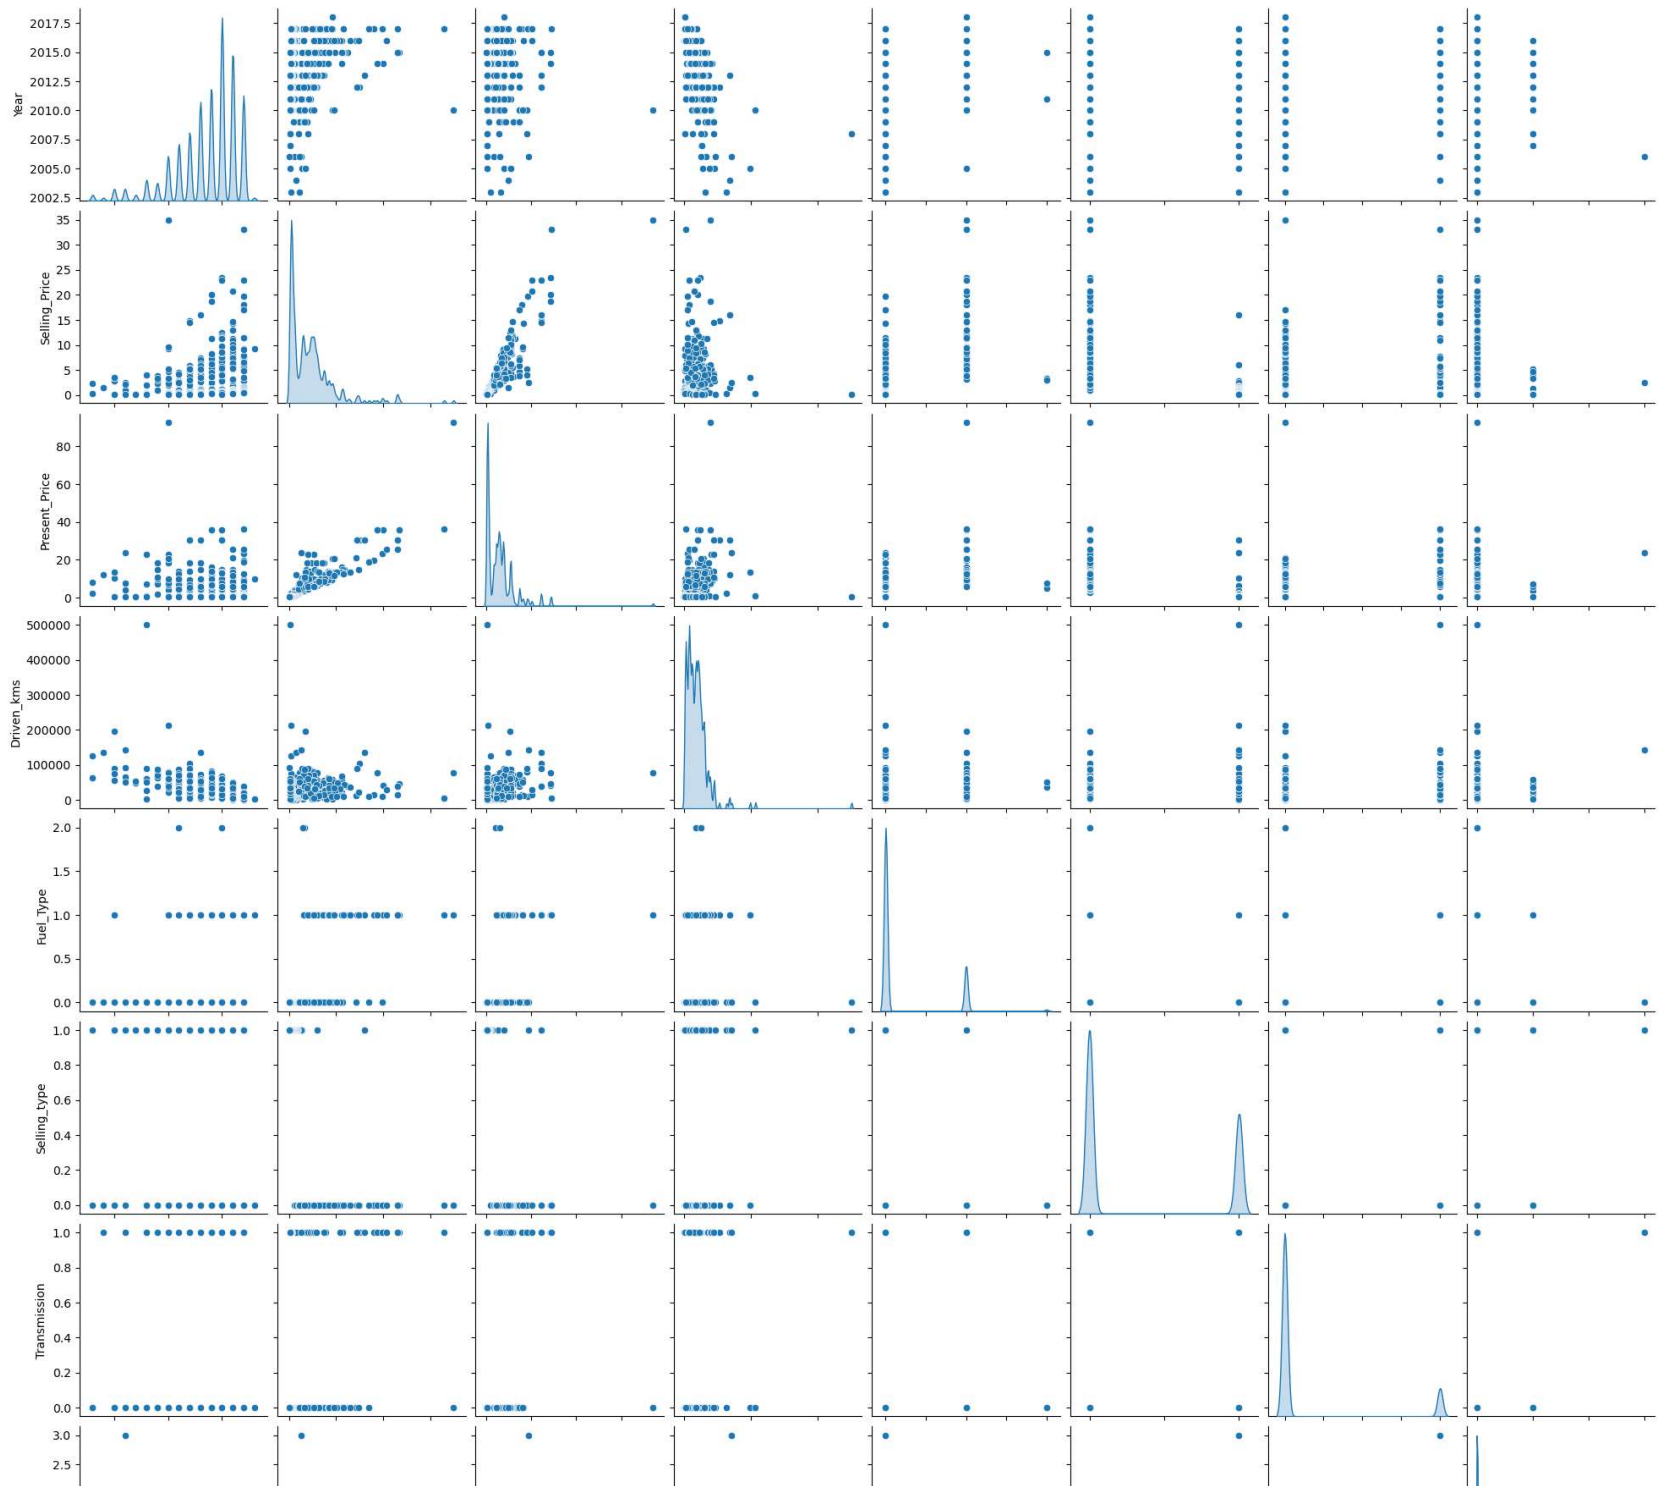
```
Petrol      239
Diesel       60
CNG          2
Name: Fuel_Type, dtype: int64
```


```
Dealer       195
Individual   106
Name: Selling_type, dtype: int64
```

```
Manual       261
Automatic     40
Name: Transmission, dtype: int64
```

```
In [63]: data.Fuel_Type.replace(regex={"Petrol":"0","Diesel":"1","CNG":"2"},inplace=True)
data.Selling_type.replace(regex={"Dealer":"0","Individual":"1"},inplace=True)
data.Transmission.replace(regex={"Manual":"0","Automatic":"1"},inplace=True)
data[["Fuel_Type","Selling_type","Transmission"]]=data[["Fuel_Type","Selling_type","Transmission"]].astype(int)
```

```
In [64]: sns.pairplot(data,diag_kind="kde", diag_kws=dict(shade=True, bw=.05, vertical=False))  
plt.show()
```



```
In [65]: y=data.Selling_Price
x=data.drop(["Selling_Price","Car_Name"],axis=1)
```

```
In [66]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2,random_state=1)
print("x train: ",x_train.shape)
print("x test: ",x_test.shape)
print("y train: ",y_train.shape)
print("y test: ",y_test.shape)
```

```
x train: (240, 7)
x test: (61, 7)
y train: (240,)
y test: (61,)
```

```
In [67]: from sklearn.metrics import r2_score
from sklearn.model_selection import cross_val_score
```

```
In [68]: cv=5
r_2 = []
CV = []
```

```
In [69]: def model(algorithm,x_train_,y_train_,x_test_,y_test_):
    algorithm.fit(x_train_,y_train_)
    predicts=algorithm.predict(x_test_)
    prediction=pd.DataFrame(predicts)
    R_2=r2_score(y_test_,prediction)
    cross_val=cross_val_score(algorithm,x_train_,y_train_,cv=cv)

    r_2.append(R_2)
    CV.append(cross_val.mean())

    # Printing results
    print(algorithm,"\n")
    print("r_2 score :",R_2,"\n")
    print("CV scores:",cross_val,"\n")
    print("CV scores mean:",cross_val.mean())

    test_index=y_test_.reset_index()["Selling_Price"]
    ax=test_index.plot(label="originals",figsize=(8,3),linewidth=2,color="r")
    ax=prediction[0].plot(label = "predictions",figsize=(12,6),linewidth=2,color="g")
    plt.legend(loc='upper right')
    plt.title("ORIGINALS VS PREDICTIONS")
    plt.xlabel("index")
    plt.ylabel("values")
    plt.show()
```




```
In [70]: from sklearn.linear_model import LinearRegression
lr = LinearRegression()
model(lr,x_train,y_train,x_test,y_test)
```

LinearRegression()

r_2 score : 0.8476231240063454

CV scores: [0.89742884 0.88694518 0.83010852 0.81465876 0.75781611]

CV scores mean: 0.8373914789815358



```
In [71]: from sklearn.linear_model import Lasso
        from sklearn.model_selection import GridSearchCV

        alphas = np.logspace(-3,3,num=14) # range for alpha

        grid = GridSearchCV(estimator=Lasso(), param_grid=dict(alpha=alphas))
        grid.fit(x_train, y_train)

        print(grid.best_score_)
        print(grid.best_estimator_.alpha)
```

0.8372790430791298

0.001


```
In [72]: ls = Lasso(alpha = grid.best_estimator_.alpha)
model(ls,x_train,y_train,x_test,y_test)
```

Lasso(alpha=0.001)

r_2 score : 0.8491511644416914

CV scores: [0.89745474 0.88725029 0.83033475 0.81478087 0.75657455]

CV scores mean: 0.8372790430791298



```
In [73]: from sklearn.linear_model import Ridge

alphas = np.logspace(-3,3,num=14)

grid2 = GridSearchCV(estimator=Ridge(), param_grid=dict(alpha=alphas))
grid2.fit(x_train, y_train)

print(grid2.best_score_)
print(grid2.best_estimator_.alpha)
```

```
0.8374514034792337
```

```
0.5878016072274912
```


```
In [74]: ridge = Ridge(alpha = 0.01)
model(ridge,x_train,y_train,x_test,y_test)
```

Ridge(alpha=0.01)

r_2 score : 0.8476639137636746

CV scores: [0.89742652 0.88696488 0.83012148 0.81467475 0.75777774]

CV scores mean: 0.8373930740726413



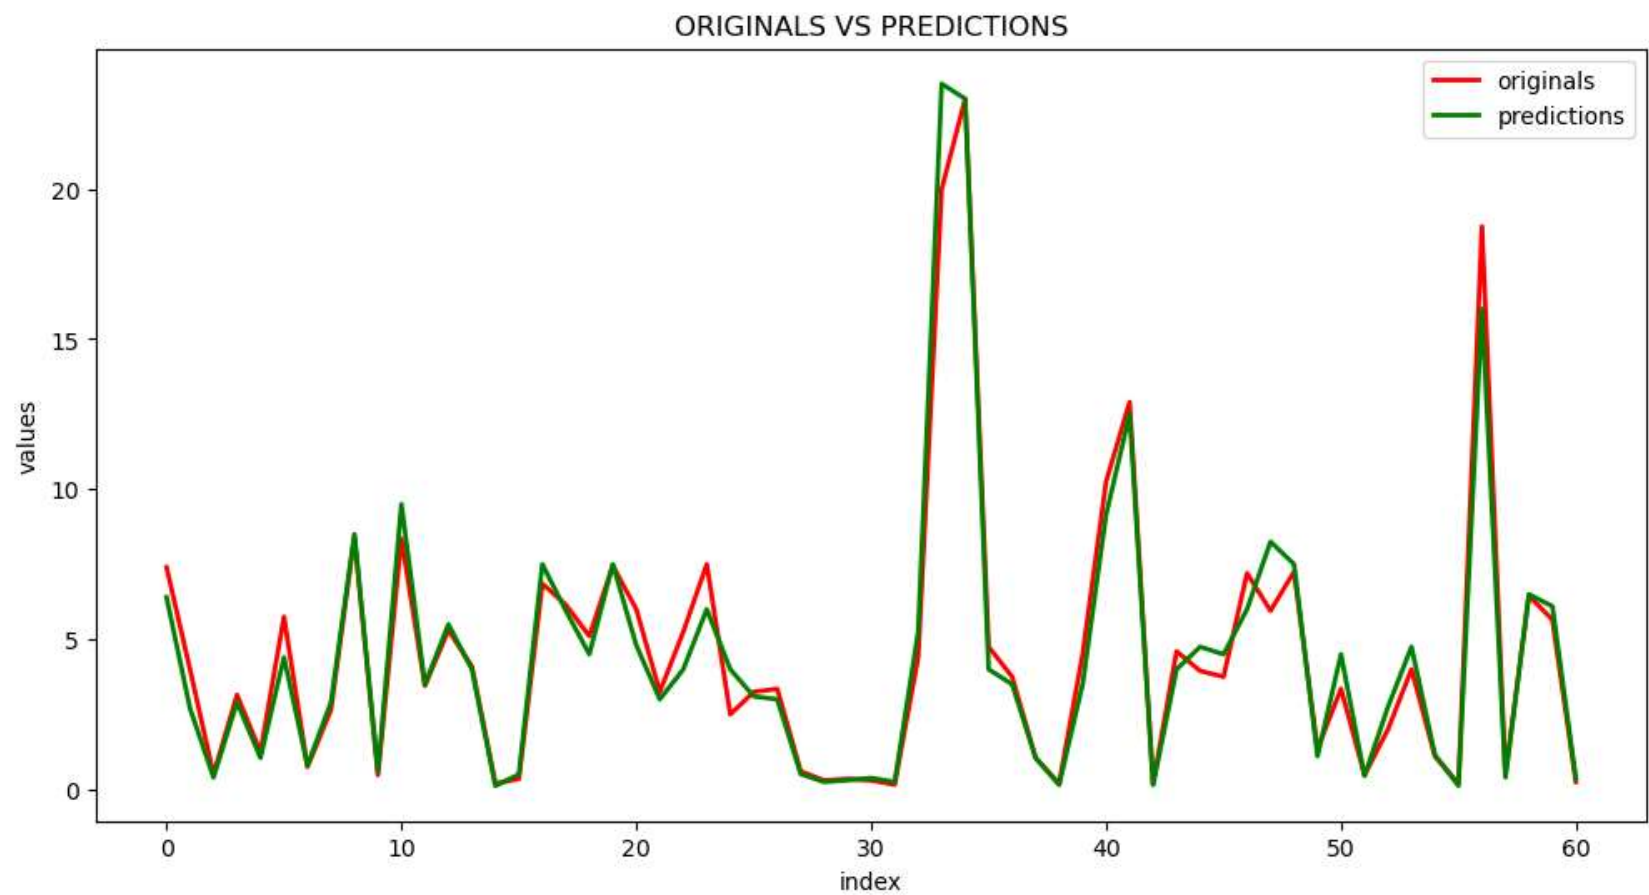
```
In [75]: from sklearn.tree import DecisionTreeRegressor  
dtr = DecisionTreeRegressor()  
model(dtr,x_train,y_train,x_test,y_test)
```

DecisionTreeRegressor()

r_2 score : 0.9626763268788968

CV scores: [0.92206829 0.84031671 0.83246236 0.90260766 0.70998591]

CV scores mean: 0.8414881860267462




```
In [76]: from sklearn.ensemble import RandomForestRegressor
rf = RandomForestRegressor(n_estimators = 100, random_state = 42)
model(rf,x_train,y_train,x_test,y_test)
```

RandomForestRegressor(random_state=42)

r_2 score : 0.9076068142402645

CV scores: [0.93398423 0.96665965 0.85051196 0.93912002 0.7208277]

CV scores mean: 0.8822207131397913



```
In [78]: Model = ["LinearRegression", "Lasso", "Ridge", "DecisionTreeRegressor", "RandomForestRegressor"]
results = pd.DataFrame({'Model': Model, 'R Squared': r_2, 'CV score mean': CV})

print(results)
```

| | Model | R Squared | CV score mean |
|---|-----------------------|-----------|---------------|
| 0 | LinearRegression | 0.847623 | 0.837391 |
| 1 | Lasso | 0.849151 | 0.837279 |
| 2 | Ridge | 0.847664 | 0.837393 |
| 3 | DecisionTreeRegressor | 0.962676 | 0.841488 |
| 4 | RandomForestRegressor | 0.907607 | 0.882221 |

In []:

In []: